

Operating Systems

Complutense University of Madrid
2020-2021

Unit 3.3: Communication and Synchronization between Threads/Processes

Contents

1 Introduction

- Classical Problems of Concurrency

2 Synchronization and communication methods

- Mutexes
- Semaphores
- Condition variables
- Shared memory between processes

3 Implementation of synchronization primitives

Contents

1 Introduction

- Classical Problems of Concurrency

2 Synchronization and communication methods

- Mutexes
- Semaphores
- Condition variables
- Shared memory between processes

3 Implementation of synchronization primitives

Concurrent processes

- **Concurrency:** simultaneous or interleaved execution of multiple instruction streams from different processes or threads
- Models
 - Multiprogramming in a single processor
 - Multiprocessors
 - Multi-computer environment (distributed computing)
- Goals
 - Share physical resources (e.g. CPUs, memory, I/O devices)
 - Share logical resources (e.g. files, data structures)
 - Accelerate calculations
 - Modularity

Classical Problems of Concurrency

- The critical section problem
- The producer-consumer problem
- The reader-writer problem
- The dining philosophers problem



The critical section problem


- Consider a concurrent program consisting of n processes/threads
- Each process/thread includes a code fragment that accesses/-modifies a shared resource
 - Critical section
- Our goal is to ensure that only one process/thread at a time can execute its critical section

The critical section problem: Example 1

```
int i;
int total_sum= 0;
for (i = 1; i <= 100; i++)
    total_sum+=i;
```



Parallelization with two threads



The critical section problem: Example 1

- Calculate $\sum_{i=1}^N i$ using multiple threads

```
int total_sum = 0; // Shared variable

void calculate_partial_sum(int ni, int nf) {
    int j = 0;
    int partial_sum = 0; // Private variable
    for (j = ni; j <= nf; j++)
        partial_sum = partial_sum + j;
    total_sum = total_sum + partial_sum;
    pthread_exit(0);
}
```

- In the event that several threads run this code concurrently, the final result may be wrong

The critical section problem: Example 1

- Potential implementation in assembly code for the critical section

```
total_sum = total_sum + partial_sum;
```

```
LDR R1,total_sum      #R1=0 (first time)
```

```
LDR R2,partial_sum    #R2=1275
```

```
ADD R1,R1,R2          #R1=1275
```

```
STR R1,total_sum      #total_sum=1275
```

The critical section problem: Example 1

- Potentially conflicting situation:

```
LDR R1,total_sum    #R1=0
LDR R2,partial_sum  #R2=1275
```

```
##### Context Switch #####
```

```
LDR R1,total_sum    #R1=0
LDR R2,partial_sum  #R2=3775
ADD R1,R1,R2        #R1=3775
STR R1,total_sum    #total_sum=3775
```

```
##### Context Switch #####
```

```
ADD R1,R1,R2        #R1=1275
STR R1,total_sum    #total_sum=1275
```

The critical section problem: Example 1

■ Solution:

- Request permission to enter the critical section
- Notify when exiting the critical section

```
int total_sum = 0; // Shared variable

void calculate_partial_sum(int ni, int nf) {
    int j = 0;
    int partial_sum = 0; // Private variable
    for (j = ni; j <= nf; j++)
        partial_sum = partial_sum + j;

    <Enter critical section>
    total_sum = total_sum + partial_sum;
    <Exit critical section>
    pthread_exit(0);
}
```

The critical section problem: Example 2

- A bank stores the balance of clients' accounts in a per-account file
 - For every money deposit in an account, the balance must be updated in the associated file

```
void perform_deposit(char *account, int amount) {  
    int balance, fd;  
    fd = open(account, O_RDWR);  
    read(fd, &balance, sizeof(int));  
    balance = balance + amount;  
    lseek(fd, 0, SEEK_SET);  
    write(fd, &balance, sizeof(int));  
    close(fd);  
    return;  
}
```

- If two processes run this code concurrently, some deposits may be performed incorrectly


The critical section problem: Example 2

```
void perform_deposit(char *account, int amount) {  
    int balance, fd;  
    fd = open(account, O_RDWR);  
    <Enter critical section>  
    read(fd, &balance, sizeof(int));  
    balance = balance + amount;  
    lseek(fd, 0, SEEK_SET);  
    write(fd, &balance, sizeof(int));  
    <Exit critical section>  
    close(fd);  
    return;  
}
```

Solution to the critical section problem


- Requirements of any solution to the critical section (CS) problem:
 - **Mutual exclusion**: only one process in the CS at a time
 - **Efficiency / Progress**: If no process is executing in the critical section, the decision on what process enters the CS depends on the processes that wish to enter the CS
 - A process should enter the CS as quickly as possible when no other process is inside the CS
 - **Avoid starvation / guarantee bounded waiting**: No process must wait forever to enter its CS
- We must also keep in mind:
 - 1 No assumptions should be made about the relative speeds of the processes or on the number of competing processes
 - 2 A process remains inside its critical section a finite amount of time

The producer-consumer problem




- The consumer will remain blocked until data is received
- The producer will block when attempting to send data over the communication mechanism when it is full

The producer-consumer problem: ring buffer



The readers-writers problem




- Readers can execute their critical section (CS) simultaneously
- If a writer is inside the CS, no other process (neither reader nor writer) can be inside the CS




The dining philosophers (Dijkstra'65)

- Five philosophers, sitting around a table, just think and eat rice:
 - They think and eat for a finite amount of time
 - The philosophers need two chopsticks to eat, which must be grabbed one after another
 - No philosopher should starve to death (avoid deadlock and starvation)
- Philosopher's code (infinite loop):
 - 1 Think...
 - 2 Grab a chopstick, grab another chopstick
 - 3 Eat
 - 4 Put the chopsticks back on the table (one after another) and go to 1.





The dining philosophers (Dijkstra'65)



The dining philosophers (Dijkstra'65)

Solutions:

- Round robin approach
 - Wastes resources
- A waiter controls access to chopsticks
 - Requires a supervisor
- Assign a number (0..4) to each chopstick and philosopher. The philosopher grabs the chopstick with a lowest number first and then the other one. After eating put them back in reverse order.
 - Bad for the last philosopher
- If a philosopher cannot grab the second chopstick, put back the first one on the table
 - What if the neighbor philosophers eat in an alternative way?

Contents

1 Introduction

- Classical Problems of Concurrency

2 Synchronization and communication methods

- Mutexes
- Semaphores
- Condition variables
- Shared memory between processes

3 Implementation of synchronization primitives

Synchronization and communication methods

- All classical problems have several things in common:
 - Processes/threads need to share information
 - They all need to access a shared variable or data structure
 - Processes/threads need to synchronize the execution
 - A process must wait for another at a given point
- We will analyze the most widely used mechanisms provided by modern OSes
 - We will study how each mechanism is used to aid in creating concurrent programs

Communication mechanisms

- Files
- Unnamed and named pipes (FIFOs)
 - Not covered by this course
- Shared memory
 - 1 Implicit: threads
 - Threads of the same process share data via global variables/data structures
 - 2 Explicit: different processes
 - Threads from different processes do not share memory
 - Processes communicate through shared memory regions
 - A specific API is needed to create shared memory regions

Synchronization mechanisms

- Operating system services:
 - Locks or mutexes
 - Semaphores
 - Condition variables
 - Signals: asynchronous, cannot be enqueued
 - Unnamed and named pipes (FIFOs)
 - Not covered by this course
- The operations supported by each mechanism must be **atomic**

Locks or mutexes

- Most suitable mechanism to solve the critical-section problem with threads (it enables to enforce MUTual EXclusion)
- A mutex has **two possible states**:
 - locked or unlocked
- It supports **two atomic operations**:
 - *Acquire the mutex (lock)*

```
lock(m) {  
    while(m.state!=unlocked)  
        .. wait ..  
  
    m.state=locked;  
}
```
 - *Release the mutex (unlock)*

```
unlock(m) { m.state=unlocked; }
```

 - it must be invoked by the thread that invoked lock() earlier
(*owner*)

Types of locks

```
lock(m) {  
    while(m.state!=unlocked)  
        .. wait ..  
  
    m.state=locked;  
}
```

3 Types of locks or mutexes

- 1 **Blocking:** The thread goes to *sleep* when waiting
 - State \Rightarrow WAITING
 - `unlock()` *wakes up* a thread \Rightarrow READY state
- 2 **Busy waiting:** Thread consumes CPU cycles while repeatedly checking the condition
 - Aka *spin lock*
 - Goal: reduce number of context switches
- 3 **Adaptive:** Thread busy waits for a while and then blocks

Locks or mutexes (*blocking*)

- A lock or mutex is a mechanism specifically tailored for synchronization between threads
 - Well-suited to the critical section problem, since it enforces Mutual EXclusion
- We can think of a mutex as an object with 3 attributes and 2 atomic methods

```

/* locked/unlocked */ lock(m) {
state_t state;         if (m->state==locked) {
                        queue_add(m->q, curThread);
                        blockCurThread();
/* queue of blocked   queue_del(m->q, curThread);
  threads */          }
queue_t q;             m->state=locked;
                        m->owner=curThread;
/* Owner thread */    }
thread_id owner;

```

```

unlock(m) {
if (m->owner==curThread) {
  m->state=unlocked;
  m->owner=NULL;
  if (m->q.notEmpty())
    wakeUpOneThread(m->q);
}
else
  error!!
}


```



Critical sections with mutexes

```
lock(m);  /* enter the critical section */  
< critical section >  
unlock(m); /* exit the critical section */
```

- The `unlock()` operation must be invoked always by the mutex's owner thread



POSIX services for mutexes

- `int pthread_mutex_init(pthread_mutex_t *mutex, pthread_mutexattr_t *attr);`
 - Initialize a mutex
- `int pthread_mutex_destroy(pthread_mutex_t *mutex);`
 - Free up the resources associated with a mutex
- `int pthread_mutex_lock(pthread_mutex_t *mutex);`
 - Acquire a mutex. If the mutex is held by another thread, the caller thread blocks until it becomes the mutex owner
- `int pthread_mutex_unlock(pthread_mutex_t *mutex);`
 - Unlocks a mutex. The caller thread must be the mutex owner.



Semaphores (Dijkstra'65)

- Synchronization mechanism
- For processes/threads in a shared-memory computer
- Object with 2 attributes
 - Queue of blocked processes/threads
 - Counter
 - initialized with a value ≥ 0
- 2 atomic operations
 - `wait()`
 - `signal()`



Operations on semaphores

```
wait(s){
    s.c = s.c - 1;
    if(s.c < 0){
        <block process>
    }
}

signal(s){
    s.c = s.c + 1;
    if (any_blocked_process){
        <Wake up a process blocked
        in wait()>
    }
}
```

Meaning of c

- $c \geq 0 \rightarrow c$ is the number of times `wait(s)` can be invoked in the concurrent program without blocking any caller process
- $c \leq 0 \rightarrow |c|$ is the number of processes blocked in the semaphore's queue

Using semaphores

■ Typical use scenarios

1 Enforce mutual exclusion

- Example: solution to the critical section problem
- Create global semaphore with $c=1$ initially
- Enclose critical sections between `wait()` and `signal()`

2 Impose synchronization restrictions associated with conditions based on integer numbers

- Example: solution to the producer/consumer problem


3 Semaphore as a general-purpose wait queue

- Examples: exercise solutions of this unit
- Create semaphore S with $c=0$ and define integer variable N to keep track of the number of processes blocked in the semaphore's queue
- N must be protected by a mutex or by another semaphore, created with $c=1$

Critical section with semaphores

- Initialize global semaphore with $c=1$

```
wait(s); /* enter the critical section */
< critical section >
signal(s); /* exit the critical section*/
```



POSIX semaphores: classification

■ 2 variants of POSIX semaphores

1 Unnamed semaphores

- Make it possible to synchronize threads or processes (usually related)
- Creation/Destruction via `sem_init()/sem_destroy()`

2 Named semaphores

- Typically used for synchronization between non-related processes
 - Each named semaphore has a global ID associated with (string)
 - Creation/Destruction via `sem_open()/sem_unlink()`
- `sem_wait()` and `sem_post()` implement the wait and signal operations, respectively
 - Can be used for either type of POSIX semaphore
- We will focus on unnamed semaphores


POSIX Semaphores: API

- `int sem_init(sem_t *sem, int shared, unsigned int val);`
 - Initialize an unnamed semaphore
 - `shared`: 0 \rightarrow threads; 1 \rightarrow processes
 - if `shared==1`, the semaphore descriptor (`sem_t`) must be stored in a memory region shared between processes
- `int sem_destroy(sem_t *sem);`
 - Free up the resources associated with an unnamed semaphore
- `int sem_wait(sem_t *sem);`
 - wait operation on the semaphore
- `int sem_post(sem_t *sem);`
 - signal operation on the semaphore

POSIX Semaphores: API

- `sem_t *sem_open(char* name,int flags,mode_t mode, unsigned int val);`
 - Opens and/or creates a named semaphore
- `int sem_close(sem_t *sem);`
 - Closes a named semaphore
- `int sem_unlink(char *name);`
 - Removes a named semaphore
- `int sem_wait(sem_t *sem);`
 - wait operation on the semaphore
- `int sem_post(sem_t *sem);`
 - signal operation on the semaphore

Producer-consumer: ring buffer



Producer-consumer with semaphores (I)

```
#define MAX_BUF          1024    /* buffer size */
#define PROD             100000  /* number of items to produce */
sem_t items;             /* number of items in the buffer */
sem_t gaps;              /* number of free gaps in the buffer */
int buffer[MAX_BUF];     /* shared buffer */

void main(void){
    pthread_t th1, th2; /* thread descriptors */
    /* semaphore initialization */
    sem_init(&items, 0, 0); sem_init(&gaps, 0, MAX_BUF);

    /* thread creation */
    pthread_create(&th1, NULL, Producer, NULL);
    pthread_create(&th2, NULL, Consumer, NULL);

    /* wait for thread completion */
    pthread_join(th1, NULL); pthread_join(th2, NULL);

    sem_destroy(&gaps); sem_destroy(&items);
    exit(0);
}
```

Producer-consumer with semaphores (II)

```
void Producer(void){
    int widx = 0; /* write index */
    int data; /* data to produce */
    int i;

    for(i=0; i < PROD; i++){
        /* produce data */
        data = generate_data();
        /* one gap less */
        sem_wait(&gaps);
        buffer[widx] = data;
        widx = (widx + 1) % MAX_BUF;
        /* added one item */
        sem_post(&items);
    }

    pthread_exit(0);
}
```

```
void Consumer(void){
    int ridx= 0; /* read index */
    int data; /* data to be consumed */
    int i;

    for(i=0; i<PROD; i++){
        /* an item will be removed */
        sem_wait(&items);
        data = buffer[ridx];
        ridx= (ridx+ 1) % MAX_BUF;
        /* one free gap more */
        sem_post(&gaps);
        do_something(data);
    }

    pthread_exit(0);
}
```



N Producers - 1 consumer with semaphores

```

#define MAX_BUF          1024    /* buffer size */
#define PROD              100000  /* number of items to produce */
#define N 2               /* Number of producers */
sem_t items;              /* number of items in the buffer */
sem_t gaps;              /* number of free gaps in the buffer */
sem_t producers;         /* To enforce mutual exclusion among producers*/
int widx=0;              /* Shared write index */
int buffer[MAX_BUF];      /* shared buffer */

void main(void){
    int i;
    pthread_t thp[N], thc; /* thread descriptors */
    /* semaphore initialization */
    sem_init(&items, 0, 0); sem_init(&gaps, 0, MAX_BUF);
    sem_init(&producers, 0, 1);

    /* thread creation */
    for (i=0;i<N;i++) pthread_create(&thp[i], NULL, Producer, NULL);
    pthread_create(&th2, NULL, Consumer, NULL);

    /* wait for thread completion */
    for (i=0;i<N;i++) pthread_join(&thp[i], NULL);
    pthread_join(th2, NULL);

    sem_destroy(&gaps); sem_destroy(&items); sem_destroy(&producers);
    exit(0);
}

```


N Producers - 1 consumer with semaphores

```
void Producer(void){
    int data; /* data to produce */
    int i;

    for(i=0; i < PROD; i++){
        /* produce data */
        data = generate_data();
        /* one gap less */
        sem_wait(&gaps);
        sem_wait(&producers);
        /* Critical section */
        buffer[widx] = data;
        widx = (widx + 1) % MAX_BUF;
        sem_post(&producers);
        /* added one item */
        sem_post(&items);
    }

    pthread_exit(0);
}
```

```
void Consumer(void){
    int ridx= 0; /* read index */
    int data; /* data to be consumed */
    int i;

    for(i=0; i<PROD; i++){
        /* an item will be removed */
        sem_wait(&items);
        data = buffer[ridx];
        ridx= (ridx+ 1) % MAX_BUF;
        /* one free gap more */
        sem_post(&gaps);
        do_something(data);
    }

    pthread_exit(0);
}
```

Readers-writers with semaphores (I)

```
int data = 5;      /* shared resource */
int nr_readers = 0; /* number of readers */
sem_t sem_nreaders; /* control access to nr_readers */
sem_t sem_read_write; /* mutual exclusion between reader-writer and
                        writer-writer */

void main(void){
    pthread_t th1, th2, th3, th4;
    sem_init(&sem_read_write, 0, 1); sem_init(&sem_nreaders, 0, 1);

    pthread_create(&th1, NULL, Reader, NULL);
    pthread_create(&th2, NULL, Writer, NULL);
    pthread_create(&th3, NULL, Reader, NULL);
    pthread_create(&th4, NULL, Writer, NULL);

    pthread_join(th1, NULL); pthread_join(th2, NULL);
    pthread_join(th3, NULL); pthread_join(th4, NULL);

    /* destroy semaphores */
    sem_destroy(&sem_read_write); sem_destroy(&sem_nreaders);
    exit(0);
}
```

Readers-writers with semaphores (II)

```
void Reader(void) {  
    while(1){  
        sem_wait(&sem_nreaders);  
        nr_readers = nr_readers + 1;  
        if (nr_readers == 1)  
            sem_wait(&sem_read_write);  
        sem_post(&sem_nreaders);  
  
        /* read data */  
        printf("%d\n", data);  
  
        sem_wait(&sem_nreaders);  
        nr_readers = nr_readers - 1;  
        if (nr_readers == 0)  
            sem_post(&sem_read_write);  
        sem_post(&sem_nreaders);  
    }  
}
```

```
void Writer(void) {  
    while(1){  
        sem_wait(&sem_read_write);  
  
        /* modify the resource */  
        data = data + 2;  
  
        sem_post(&sem_read_write);  
    }  
}
```

Condition variables

- Synchronization mechanism for threads
- Each condition variable has a wait queue and a mutex associated with it
 - The mutex is typically shared between multiple condition variables in the same concurrent program
- Condition variables support three operations:
 - 1 `cond_wait()`: the caller thread blocks in the wait queue
 - 2 `cond_signal()`: wake up a thread waiting in the condition variable's wait queue (if there are blocked threads)
 - 3 `cond_broadcast()`: wake up all threads waiting in the condition variable's wait queue (if there are blocked threads)
- These operations must be invoked in a code snippet between `lock(mutex)` and `unlock(mutex)`

Operations on condition variables (I)

```
void cond_wait(lock_t m, vc_t varC ) {  
    queue_add(varC->queue, curThread);  
    unlock(m);  
    park(); // the thread goes to sleep  
    lock(m);  
}
```

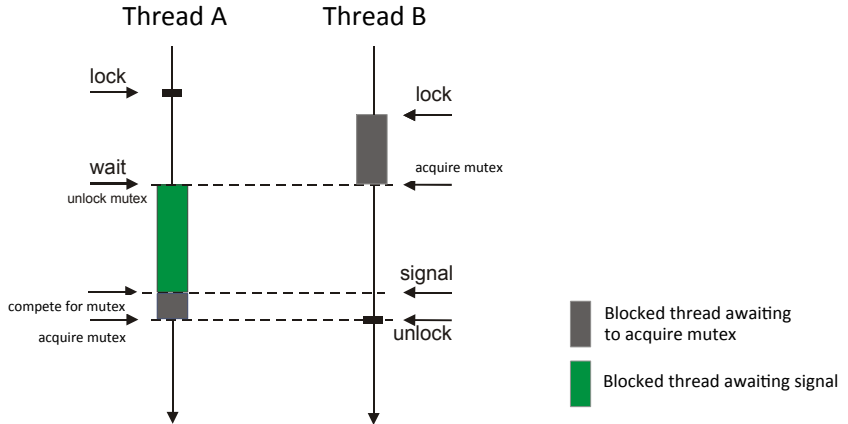
- The caller thread must be the owner of the mutex
- `cond_wait()` **always** blocks the caller thread
 - Before going to sleep, the thread unlocks the mutex so that another thread can acquire it
- When the thread is awoken, it attempts to reacquire (lock) the mutex again (the thread may block)
- When the `cond_wait()` operation returns, the caller thread is the owner of the mutex again

Operations on condition variables (II)

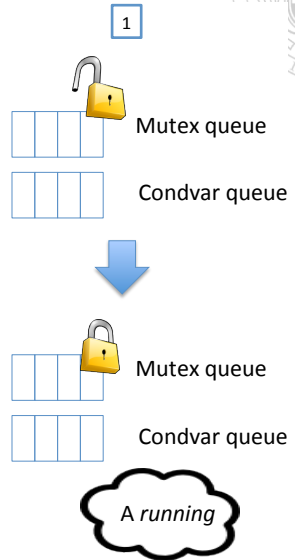

```
/* Wakes up one thread from the wait queue */  
void cond_signal (vc_t varC ) {  
    if (!isEmpty(varC->queue))  
        unpark(queue_remove(varC->queue))  
}  
  
/* Wakes up all threads in the wait queue */  
void cond_broadcast (vc_t varC ) {  
    while (!isEmpty(varC->queue))  
        unpark(queue_remove(varC->queue))  
}
```

- It is strongly advisable that these operations are invoked in a code snippet between `lock(mutex)` and `unlock(mutex)`

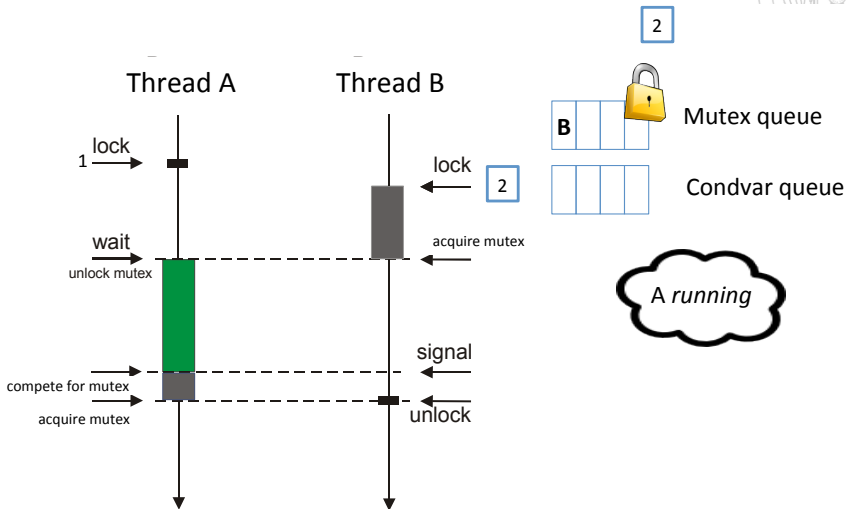
Operations on condition variables (III)



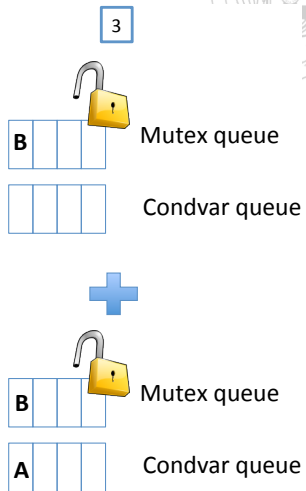
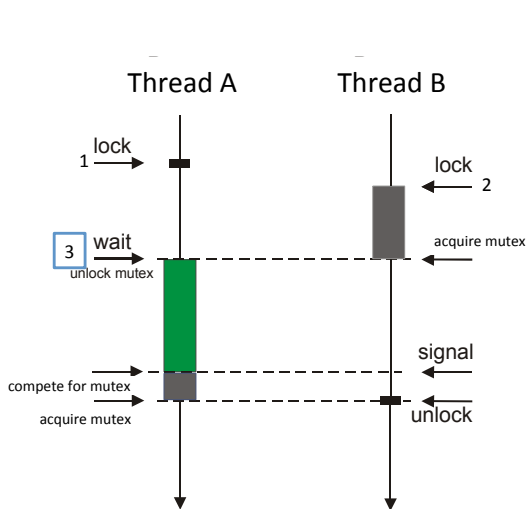
Using condition variables



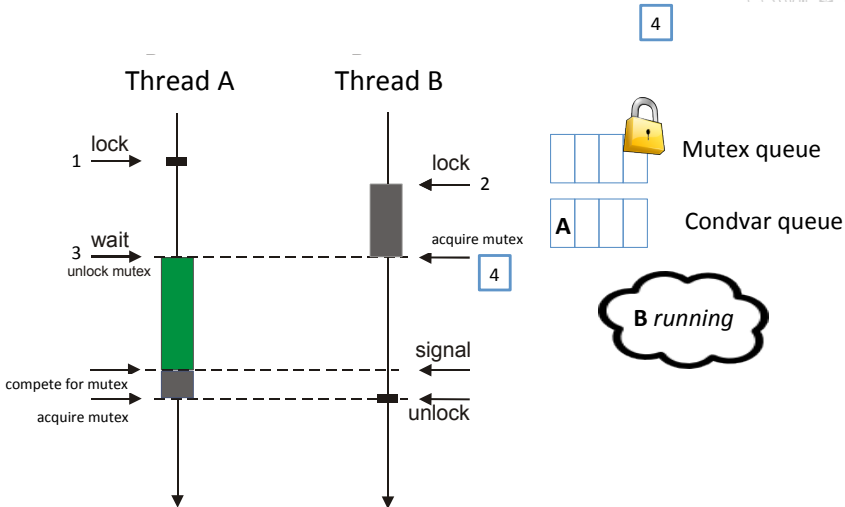
Using condition variables



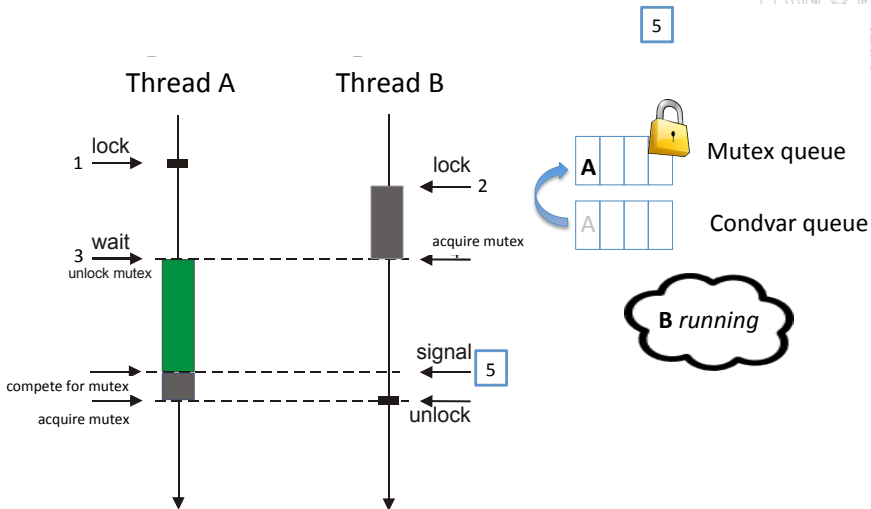
Using condition variables



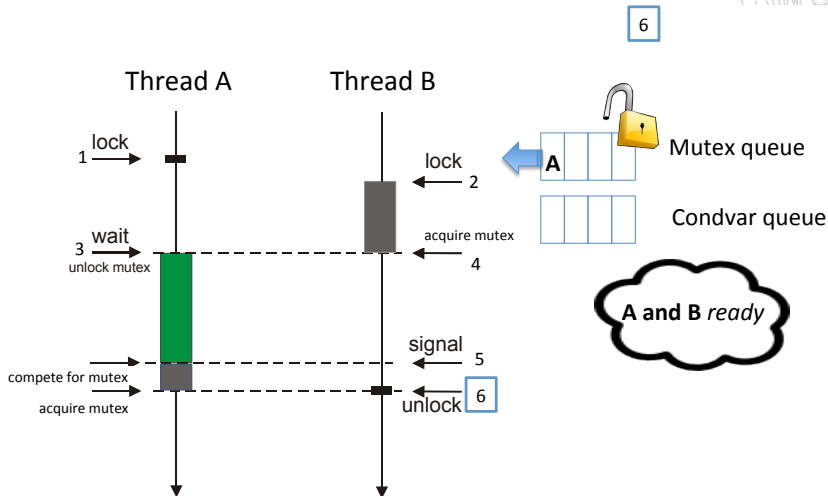
Using condition variables



Using condition variables



Using condition variables



Using condition variables with a mutex

■ Thread A

```
lock(mutex); /* enter the CS */
<operations on the shared resource (mutual exclusion)>
while (condition related to the resource == false)
    cond_wait(condition, mutex); /* thread blocks */
<desired actions once the condition is met>
unlock(mutex); /* exit the CS */
```

■ Thread B

```
lock(mutex); /* enter the CS */
<operations on the shared resource (mutual exclusion)>
/* Since we may have affected the condition, wake up other thread */
cond_signal(condition, mutex);
<more operations on the shared resource>
unlock(mutex); /* exit the CS */
```

We must use `while` to re-evaluate the condition

Producer-Consumer with cond. variables (I)

- For simplicity, assume that the ring buffer is already implemented as an abstract data type (`cbuffer_t`)
 - Operations on `cbuffer_t`:
 - `boolean is_empty(cbuffer_t cb);`
 - `boolean is_full(cbuffer_t cb);`
 - `void add(cbuffer_t cb, item_t data);`
 - `item_t remove(cbuffer_t cb);`
- Operation restrictions
 - 1 Thread-unsafe implementation
 - The operations cannot be invoked simultaneously from multiple threads
 - A lock is needed to serialize accesses to the data structure
 - 2 `add()` cannot be invoked if buffer is full
 - 3 `remove()` cannot be invoked if buffer is empty

Producer-Consumer with cond. variables (II)

Global variables

```
cbuffer_t b; /* shared ring buffer (max capacity: N items) */  
mutex m; /* Mutual exclusion when accessing buffer */  
condvar prod, cons; /* To block producer/consumer, respectively */
```

```
void producer() {  
    item_t data;  
  
    while (true) {  
        data=produce();  
        lock(m);  
  
        add(b,data);  
  
        unlock(m);  
        delay(...);  
    }  
}
```

```
void consumer() {  
    item_t data;  
  
    while (true) {  
        lock(m);  
  
        data=remove(b);  
  
        unlock(m);  
        do_something(data);  
    }  
}
```




Producer-Consumer with cond. variables (II)

Global variables

```
cbuffer_t b; /* shared ring buffer (max capacity: N items) */
mutex m; /* Mutual exclusion when accessing buffer */
condvar prod, cons; /* To block producer/consumer, respectively */
```

```
void producer() {
    item_t data;

    while (true) {
        data=produce();
        lock(m);

        while (is_full(b))
            cond_wait(prod,m);

        add(b,data);

        unlock(m);
        delay(...);
    }
}
```

```
void consumer() {
    item_t data;

    while (true) {
        lock(m);

        while (is_empty(b))
            cond_wait(cons,m);

        data=remove(b);

        unlock(m);
        do_something(data);
    }
}
```



Producer-Consumer with cond. variables (II)

Global variables

```
cbuffer_t b; /* shared ring buffer (max capacity: N items) */
mutex m; /* Mutual exclusion when accessing buffer */
condvar prod, cons; /* To block producer/consumer, respectively */
```

```
void producer() {
    item_t data;

    while (true) {
        data=produce();
        lock(m);

        while (is_full(b))
            cond_wait(prod,m);

        add(b,data);

        cond_signal(cons);

        unlock(m);
        delay(...);
    }
}
```

```
void consumer() {
    item_t data;

    while (true) {
        lock(m);

        while (is_empty(b))
            cond_wait(cons,m);

        data=remove(b);

        cond_signal(prod);

        unlock(m);
        do_something(data);
    }
}
```

POSIX Services (II)


- `int pthread_cond_init(pthread_cond_t* cond, pthread_condattr_t* attr);`
 - Initialize a condition variable
- `int pthread_cond_destroy(pthread_cond_t *cond);`
 - Free up the resources associated with a condition variable
- `int pthread_cond_signal(pthread_cond_t *cond);`
 - Unblock one thread waiting in the queue associated with a condition variable
 - It has no effect if the wait queue is empty (different behavior to that of the signal operation for semaphores).
- `int pthread_cond_broadcast(pthread_cond_t *cond);`
 - Unblock all threads waiting in the queue associated with a condition variable
- `int pthread_cond_wait(pthread_cond_t* cond, pthread_mutex_t* mutex);`
 - Blocks the caller thread until it is eventually awoken by `pthread_cond_signal()` or `pthread_cond_broadcast()`

Monitors

- Semaphores enable us to:
 - 1 Enforce mutual exclusion
 - 2 Impose synchronization restrictions in the program
- However, implementations with semaphores are usually hard to follow and difficult to debug
- Monitors make it possible to separate both aspects by
 - using mutexes to ensure mutual exclusion
 - and condition variables to impose more elaborate synchronization restrictions based on conditions depending on the state of shared variables among threads
- More suitable for object oriented programming (e.g. Java)
- A monitor is not a synchronization mechanism provided by the OS but instead we must build it specifically for our concurrent program

Monitors

- A monitor is an object consisting of 3 components:
 - 1 Monitor procedures (*methods*)
 - 2 Variables and data structures shared between threads (*attributes*)
 - 3 A mutex + N condition variables
 - The mutex is implicit in some monitor implementations
- Monitor procedures (methods) can be invoked concurrently by multiple threads
 - Concurrency issues are addressed inside each monitor procedure



Monitors: Implementation (I)

- A thread “enters the monitor” when it invokes a monitor procedure; the thread “exits the monitor” when it completes the execution of the procedure
- Implementation restrictions:
 - 1 Only one active (non-blocked) thread can be inside a monitor procedure at a time
 - The thread running inside the monitor accesses the monitor’s variables in mutual exclusion with respect to other threads
 - 2 One or several threads can be blocked inside the monitor
 - Threads may remain blocked in the condition variables associated with the monitor

Monitors: Implementation (II)

- In the C/POSIX-threads monitor implementation, the aforementioned restrictions can be imposed as follows:
 - A monitor procedure is a C function with a single return point
 - The body of the monitor procedure is enclosed between `lock()` and `unlock()` of the monitor mutex

```
monitor_procedure() {  
    pthread_mutex_lock(&mutex); /* enter the monitor */  
    <procedure body>  
    pthread_mutex_unlock(&mutex); /* exit the monitor */  
}
```

- Every thread blocked in a monitor's condition variable, releases the mutex implicitly when invoking `cond_wait()`

Using monitors

- Threads in a concurrent program invoke the monitor procedures for different reasons:
 - 1 Modify/access the monitor variables safely
 - 2 Enter/exit a critical section

```
entry_procedure();  
<< critical section >>  
exit_procedure();
```


Producer-consumer with monitor

```
#define MAX_BUFFER 1024 /* buffer size */
#define DATA_TO_PRODUCE 10000
pthread_mutex_t mutex; /* monitor's mutex */
pthread_cond_t c_full; /* to block the producer */
pthread_cond_t c_empty; /* to block the consumer */
int nr_items; /* # of items in the buffer */
int buffer[MAX_BUFFER]; /* shared buffer */
int ridx, widx; /* R/W positions in the buffer */

void init_monitor(void); /* Initialize monitor */
void destroy_monitor(void); /* Free up monitor resources */
/* Monitor procedures */
void produce(int item); /* Insert item into the buffer */
int consume(void); /* Extract an item from the buffer */
```

Producer-consumer with monitor

```
void Producer(void) {  
    int i,item;  
    for (i=0;i<DATA_TO_PRODUCE;i++){  
        item=... generate an item ...  
        produce(item);  
    }  
    pthread_exit(0);  
}
```

```
void Consumer(void) {  
    int i,item;  
    for (i=0;i<DATA_TO_PRODUCE;i++){  
        item=consume();  
        ... Do something with item ...  
    }  
    pthread_exit(0);  
}
```

```
int main(int argc, char *argv[]){  
    pthread_t th1, th2;  
    init_monitor();  
    pthread_create(&th1, NULL, Producer, NULL);  
    pthread_create(&th2, NULL, Consumer, NULL);  
    pthread_join(th1, NULL); pthread_join(th2, NULL);  
    destroy_monitor();  
    return 0;  
}
```

Producer-consumer with monitor

```
void init_monitor(void)
{
    pthread_cond_init(&c_full, NULL);
    pthread_cond_init(&c_empty, NULL);
    pthread_mutex_init(&mutex, NULL);
    ridx=widx=nr_items=0;
}

void destroy_monitor(void)
{
    pthread_mutex_destroy(&mutex);
    pthread_cond_destroy(&c_full);
    pthread_cond_destroy(&c_empty);
}
```

Producer-consumer with monitor

```
void produce(int item) {  
    /* "enter the monitor" */  
    pthread_mutex_lock(&mutex);  
  
    /* block while buffer full */  
    while (nr_items == MAX_BUFFER)  
        pthread_cond_wait(&c_full,&mutex);  
  
    buffer[widx] = item;  
    widx= (widx+ 1) % MAX_BUFFER;  
    nr_items++;  
  
    /* buffer is not empty */  
    pthread_cond_signal(&c_empty);  
  
    /* "exit the monitor" */  
    pthread_mutex_unlock(&mutex);  
}
```

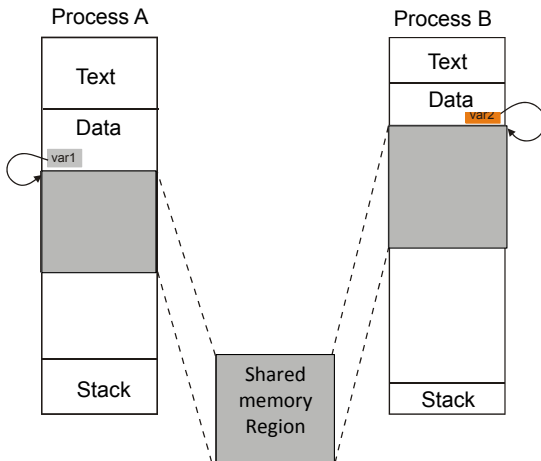
```
int consume(void) {  
    int item;  
    /* "enter the monitor" */  
    pthread_mutex_lock(&mutex);  
  
    /* block while buffer empty */  
    while (nr_items == 0)  
        pthread_cond_wait(&c_empty,&mutex);  
  
    item = buffer[ridx];  
    ridx= (ridx + 1) % MAX_BUFFER;  
    nr_items--;  
  
    /* buffer is not full */  
    pthread_cond_signal(&c_full);  
  
    /* "exit the monitor" */  
    pthread_mutex_unlock(&mutex);  
    return item;  
}
```

Communication mechanisms

- Files
- Unnamed and named pipes (FIFOs)
 - Not covered by this course
- Shared memory
 - 1 Implicit: threads
 - Threads of the same process share data via global variables/data structures
 - 2 Explicit: different processes
 - Threads from different processes do not share memory
 - Processes communicate through shared memory regions
 - A specific API is needed to create shared memory regions

Shared memory between processes

- Independent declaration of variables inside each process that refers to the same physical memory location



POSIX shared memory

- `void *mmap(void *addr, size_t length, int prot, int flags, int fd, off_t offset);`
 - maps part of a file (described by `fd`) into memory and returns a pointer to that memory region (`addr`)
 - The memory region can be either shared or private:
 - flags: `MAP_SHARED` or `MAP_PRIVATE`
 - The memory region can be declared without an associated file:
 - flags: `MAP_ANONYMOUS` (shared between parent and child)
 - By using `shm_open` to obtain a file descriptor
- `int munmap(void *addr, size_t length);`
 - delete the mappings for the specified address range, causing further references to addresses within the range to generate invalid memory references
- `int msync(void *addr, size_t len, int flags);`
 - writes modified whole pages back to the filesystem and updates the file modification time.

Summary

- Threads:
 - Shared memory (global variables)
 - Mutexes, condition variables, semaphores, monitors.
- Related processes (`fork()`):
 - Unnamed Pipes, shared memory
 - Unnamed semaphores (memory-mapped)
- Non-related processes:
 - Named semaphores
 - Named Pipes, shared memory

P&C with shared memory and semaphores

■ Producer:

- Creates named semaphores (`sem_open`)
- Creates a file (`open`)
- Assigns space to it (`ftruncate`)
- Maps the file into the process address space (`mmap`)
- Uses the shared memory region
- Unmaps the shared memory region (`munmap`)
- Closes the file and deletes it

■ Consumer:

- Opens the named semaphores (`sem_open`)
- The consumer must wait until the file is created to open it (`open`)
- Maps the file into the process address space (`mmap`)
- Uses the shared memory region
- Closes the file

Producer's code

```
#define MAX_BUFFER      1024    /* buffer size */
#define DATA_TO_PRODUCE 100000 /* # elements to produce */
sem_t *elements; /* # of elements in the buffer */
sem_t *gaps; /* # of free gaps in the buffer */

void main(int argc, char *argv[]){
    int shd;
    int *buffer; /* shared buffer */

    /* the producer creates the file */
    shd = open("BUFFER", O_CREAT|O_WRONLY, 0700);
    ftruncate(shd, MAX_BUFFER * sizeof(int));

    /* Maps the file into the process address space */
    buffer = (int*) mmap(NULL, MAX_BUFFER * sizeof(int),
        PROT_WRITE, MAP_SHARED, shd, 0);
```

Producer's code (II)

```
/* The producer creates the semaphores */
elements = sem_open("ELEMENTS", O_CREAT, 0700, 0);
gaps = sem_open("GAPS", O_CREAT, 0700, MAX_BUFFER);

/* core producer's code */
Producer(buffer);

/* Unmap shared buffer */
munmap(buffer, MAX_BUFFER * sizeof(int));
close(shd); /* close the shared memory region */
unlink("BUFFER"); /* delete the shared memory region */

sem_close(elements);
sem_close(gaps);
sem_unlink("ELEMENTS");
sem_unlink("GAPS");
}
```

Consumer's code

```
#define MAX_BUFFER      1024    /* buffer size */
#define DATA_TO_PRODUCE 100000 /* # elements to produce */
sem_t *elements; /* # of elements in the buffer */
sem_t *gaps; /* # of free gaps in the buffer */

void main(int argc, char *argv[]){
    int shd;
    int *buffer; /* shared buffer */

    /* the consumer opens the file */
    shd = open("BUFFER", O_RDONLY);

    /* Maps the file into the process address space */
    buffer = (int *) mmap(NULL, MAX_BUFFER * sizeof(int),
                          PROT_READ, MAP_SHARED, shd, 0);
```

Consumer's code (II)

```
/* Consumer opens semaphores */
elementos = sem_open("ELEMENTS", 0);
huecos    = sem_open("GAPS", 0);

/* consumer's core processing */
Consumer(buffer);

/* unmap shared buffer */
munmap(buffer, MAX_BUFFER * sizeof(int));
close(shd); /* close shared memory object */

/* close semaphores */
sem_close(elementos);
sem_close(gaps);
}
```

Producer's core processing

```
void Producer(int *buffer)
{
    int pos = 0; /* write index */
    int item;    /* data to produce */
    int i;

    for(i=0; i < DATA_TO_PRODUCE; i++ ) {
        item = produce_item();
        sem_wait(gaps);
        buffer[pos] = item;
        pos = (pos + 1) % MAX_BUFFER;
        sem_post(elements);
    }
}
```

Consumer's core processing

```
void Consumer(char *buffer)
{
    int pos = 0; /* read index */
    int i, item;

    for(i=0; i < DATA_TO_PRODUCE; i++ ) {
        sem_wait(elements);
        dato = buffer[pos];
        pos = (pos + 1) % MAX_BUFFER;
        sem_post(gaps);
        printf("Received %d\n", item);
    }
}
```

Contents

1 Introduction

- Classical Problems of Concurrency

2 Synchronization and communication methods

- Mutexes
- Semaphores
- Condition variables
- Shared memory between processes

3 Implementation of synchronization primitives

Implementation of lock()/unlock()

- The OS ensures that lock() and unlock() are **atomic operations**
 - Note that the implementation entails solving the critical section problem (shared variables)

```
lock(m) {  
  if (m->state==locked) {  
    queue_add(m->q, curThread);  
    blockCurThread();  
    queue_del(m->q, curThread);  
  }  
  m->state=locked;  
  m->owner=curThread;  
}
```

```
unlock(m) {  
  if (m->owner==curThread) {  
    m->state=unlocked;  
    m->owner=NULL;  
    if (m->q.notEmpty())  
      wakeUpOneThread(m->q);  
  }  
  else  
    error!!  
}
```

Solutions to the critical section problem

Types of solutions

- Based on busy waiting
 - Without HW support
 - Based on control variables (Peterson 1981)
 - With HW support
 - Test And Set (TAS), XCHG, LL/SC
- Without busy waiting (blocked waiting)
 - OS support needed
 - The OS blocks the process/thread

Machine instructions

- A machine instruction is used to update the contents of a memory address atomically
- Solution works for *any* number of processes/threads:
 - **RMW** memory cycle (*read/modify/write*)
- These special instructions are not negatively affected by other instructions
- The mechanism can be applied to multiple critical sections
- Simple mechanism that can be easily verified

Machine instructions: example

General

- Test and set (T&S)
- Fetch and add (F&A)
- Swap/Exchange
- Compare and Swap (exchange)
- Load link/ Store conditional (LL/SC)

Intel (x86)

- Many atomic instructions supported
- F&A lock; `xaddl eax, [mem_addr];`
- XCHG `xchg eax, [mem_addr]`
- CMPXCHG -> `lock cmpxchg [mem_addr], eax`

ARM (and others)

- LL/SC LDREX y STREX

Semantics of Swap/Exchange

Exchange

```
xchg src, dst
```

```
-----
```

```
rtmp  <- Mem [src]  
Mem [src] <- Mem [dst]  
Mem [dst] <- rtmp
```

- It is a machine instruction rather than a function
 - It is atomic, uninterruptible
- Swaps two values (Both potentially in memory)
 - On Intel processors, just one of both (src or dst) can be on memory

Usage of Swap/Exchange

- Solution to the **critical section problem** with **XCHG**

Implementation

```
/* it may be stored in a register */  
tmp = 1;  
/* Busy wait */  
while( tmp== 1)  
    xchg(MemAddr, tmp);  
  
Critical_section();  
  
*MemAddr= 0;
```

Semantics of LL/SC

Load Link

```
ll src
-----
    rout <- Mem [src]
```

Store Conditional

```
sc src, value
-----
    if nobody accessed src since the last LL
        Mem[src]= value
        rout <- 1
    sino
        rout <- 0
```

- LL/SC constitute two independent machine instructions
 - LL takes care of loading the data in a register
 - SC only stores the data if no writes were performed to that memory location prior to the execution of LL

Usage of LL/SC

- Solution to the **critical section problem** with LL/SC

Implementation

```
while (1) {  
    while(ll(MemAddr)== 1);  
    if (sc(MemAddr,1)==1) break;  
    /* otherwise, execute Load-Link again */  
}  
Critical_section();  
*MemAddr= 0;
```